

Stochastic Machine Learning

01 - Introduction

Thorsten Schmidt

Abteilung für Mathematische Stochastik

www.stochastik.uni-freiburg.de
thorsten.schmidt@stochastik.uni-freiburg.de

WS 2020/21

Introduction → Machine Learning Basics

Definition

A computer program learns from experience E with respect to tasks T , if its performance P improves with experience E .

This quite vague definition allows us to develop some intuition about the situation.

- ▶ **Experience** is given by an increasing sequence of observations, for example X_1, X_2, \dots, X_t could represent the information at time t . This is typically decoded in a **filtration**: a filtration is an increasing sequence of sub- σ -fields $(\mathcal{F}_t)_{t \in \mathcal{T}}$.
- ▶ The performance is often measured in terms of an **utility function**. For example the utility at time t could be given by $U(X_t)$ with an function U . U could of course depend on more variables. One could also look for the accumulated utility

$$\sum_{t=1}^T U(X_t).$$

One very simple learning algorithm is linear regression, a classical statistical concept. Here it arises as an example of **supervised learning**.

Example (Linear Regression)

Suppose we observe pairs $(x_i, y_i)_{i=1, \dots, n}$ and want to predict y on basis of x . **Linear regression** requires

$$\hat{y}(x) = \beta x$$

with some weight $\beta \in \mathbb{R}$. We specify a loss function¹¹

$$\text{RSS}(\beta) := \sum_{i=1}^n (y_i - \hat{y}(x_i))^2$$

and minimize over β .

One could choose $-\text{MSE}$ as utility function. So how does the system **learn**?

¹¹Given by the Residual Sum of Squares here.

The system learns by maximizing the utility, i.e. minimizing the MSE for each n . And additional data will lead to a better prediction. We will later see that this is in a certain sense indeed optimal.

We use the **first-order condition** to derive the solution letting $\mathbf{x} = (x_1, \dots, x_n)$ and similar for \mathbf{y} ,

$$\begin{aligned} 0 &= \partial_{\beta} (\mathbf{y} - \beta \mathbf{x})^2 = \partial_{\beta} (\mathbf{y}^2 - 2\mathbf{y}^{\top} \beta \mathbf{x} + \beta^2 \mathbf{x}^{\top} \mathbf{x}) \\ \Leftrightarrow \quad 0 &= -2\mathbf{x}^{\top} \mathbf{y} + 2\beta \mathbf{x}^{\top} \mathbf{x} \end{aligned}$$

such that we obtain

$$\hat{\beta} = (\mathbf{x}^{\top} \mathbf{x})^{-1} \mathbf{x}^{\top} \mathbf{y}.$$

Note that typically one considers affine functions of x without mentioning, i.e. one looks at functions $y = \alpha + \beta x$. This can simply be achieved with the linear approach by augmenting \mathbf{x} by an additional entry 1.

- ▶ Of course many generalizations are possible:
- ▶ To higher dimensions: consider data vectors $(\mathbf{x}_i, \mathbf{y}_i)$, $i = 1, \dots, n$,
- ▶ To nonlinear functions: include x_i^1, \dots, x_i^p into the covariates
- ▶ and many more.

Let us consider a linear regression in python.

```
import yfinance as yf
import matplotlib.pyplot as plt

DAX = yf.Ticker('%5Egdaxi')
DAX_History = DAX.history(start="2020-01-01", end="2020-10-26")


plt.figure(figsize=(10,10))
plt.plot(DAX_History.index, DAX_History['Close'])

# Linear Regression example: regress tomorrow on today
x = DAX_History['Close'][:-1] # without the last value
y = DAX_History['Close'][1:] # without the first value

import numpy as np
from numpy import array
from sklearn.linear_model import LinearRegression

model = LinearRegression()
x = array(x).reshape(-1,1) # The lin
y = array(y).reshape(-1,1)
model.fit(x, y) # values in model.int

# Give a very sophisticated plot
import seaborn as sns; sns.set_theme()
ax = sns.regplot(x=x, y=y)
plt.show()
```



Could we improve this ? Suggestions ?

Let us consider a linear regression in python.

```
import yfinance as yf
import matplotlib.pyplot as plt

DAX = yf.Ticker('%5Egdaxi')
DAX_History = DAX.history(start="2020-01-01", end="2020-10-26")


plt.figure(figsize=(10,10))
plt.plot(DAX_History.index, DAX_History['Close'])

# Linear Regression example: regress tomorrow on today
x = DAX_History['Close'][:-1] # without the last value
y = DAX_History['Close'][1:] # without the first value

import numpy as np
from numpy import array
from sklearn.linear_model import LinearRegression

model = LinearRegression()
x = array(x).reshape(-1,1) # The linear model
y = array(y).reshape(-1,1)
model.fit(x, y) # values in model.int

# Give a very sophisticated plot
import seaborn as sns; sns.set_theme(context='notebook')
ax = sns.regplot(x=x, y=y)
plt.show()
```



Could we improve this ? Suggestions ?

What is the difference to Statistics ?

In a statistical approach we start with a **parametric model**:

$$Y_i = \alpha + \beta x_i + \epsilon_i, \quad i = 1, \dots, n$$

and assume that $\epsilon_1, \dots, \epsilon_n$ have a certain structure (for example, i.i.d. and $\mathcal{N}(0, \sigma^2)$). The one can derive (see, e.g. Czado & Schmidt (2011)) **optimal estimators** for α and β . One can also relax the assumptions and gets weaker results.

So what ? What are the advantages of the statistical approach ?

One particular outcome is that we are able to provide **confidence intervals**, **predictive intervals** and **test** hypothesises.

Questions

- ▶ What is the definition of Machine Learning?
- ▶ Give examples
- ▶ Give surprising examples
- ▶ Derive the main equation of linear regression
- ▶ (do it in 1 dimension first - this goes back to Gauss)
- ▶ Write your own python code, providing a linear regression on your favourite stock
- ▶ Do this with your least favourite stock
- ▶ Can you regress two stocks on each other ?
- ▶ Can you predict better the value of the stock tomorrow ? (You can also research on this ...)